NAME

    magic - file command's magic pattern file

DESCRIPTION

    This manual page documents the format of the magic file as used by the file(1) command. file(1) identifies the type of a file using, among other tests, a test for whether the file contains certain "magic patterns". The file */etc/magic* specifies what magic numbers are to be tested for, what message to print if a particular magic number is found, and additional information to extract from the file.

    Each line of the file specifies a test to be performed. A test compares the data starting at a particular offset in the file with a byte value, a string or a numeric value. If the test succeeds, a message is printed. The line consists of the following fields:

offset   A number specifying the offset, in bytes, into the file of the data which is to be tested.

type   The type of the data to be tested. The possible values are:

        byte   A one-byte value.

        short   A two-byte value in this machine's native byte order.

        long   A four-byte value in this machine's native byte order.

        quad   An eight-byte value in this machine's native byte order.

        float   A 32-bit single precision IEEE floating point number in this machine's native byte order.

        double   A 64-bit double precision IEEE floating point number in this machine's native byte order.

        string   A string of bytes. The string type specification can be optionally followed by /[Bbc]*. The "B" flag compacts whitespace in the target, which must contain at least one whitespace character. If the magic has n consecutive blanks, the target needs at least n consecutive blanks to match. The "b" flag treats every blank in the target as an optional blank. Finally the "c" flag, specifies case insensitive matching: lowercase characters in the magic match both lower and upper case characters in the target, whereas upper case characters in the magic only match uppercase characters in the target.

pstring   A Pascal-style string where the first byte is interpreted as the an unsigned length. The string is not NUL terminated.

date   A four-byte value interpreted as a UNIX date.

qdate   An eight-byte value interpreted as a UNIX date.

ldate   A four-byte value interpreted as a UNIX-style date, but interpreted as local time rather than UTC.

qldate   An eight-byte value interpreted as a UNIX-style date, but interpreted as local time rather than UTC.

beshort   A two-byte value in big-endian byte order.

belong   A four-byte value in big-endian byte order.

bequad   An eight-byte value in big-endian byte order.

befloat   A 32-bit single precision IEEE floating point number in big-endian byte order.

bedouble   A 64-bit double precision IEEE floating point number in big-endian byte order.

bedate   A four-byte value in big-endian byte order, interpreted as a Unix date.

beqdate   An eight-byte value in big-endian byte order, interpreted as a Unix date.

beldate   A four-byte value in big-endian byte order, interpreted as a UNIX-style date, but interpreted as local time rather than UTC.

beqldate   An eight-byte value in big-endian byte order, interpreted as a UNIX-style date, but interpreted as local time rather than UTC.

bestring16   A two-byte unicode (UCS16) string in big-endian byte order.

leshort   A two-byte value in little-endian byte order.

lelong   A four-byte value in little-endian byte order.

lequad   An eight-byte value in little-endian byte order.

lefloat   A 32-bit single precision IEEE floating point number in little-endian byte
                order.

ledouble   A 64-bit double precision IEEE floating point number in little-endian byte
                order.

ledate   A four-byte value in little-endian byte order, interpreted as a UNIX date.

leqdate   An eight-byte value in little-endian byte order, interpreted as a UNIX date.

leldate   A four-byte value in little-endian byte order, interpreted as a UNIX-style
                date, but interpreted as local time rather than UTC.

leqldate   An eight-byte value in little-endian byte order, interpreted as a UNIX-style
                date, but interpreted as local time rather than UTC.

lestring16   A two-byte unicode (UCS16) string in little-endian byte order.

melong   A four-byte value in middle-endian (PDP-11) byte order.

medate   A four-byte value in middle-endian (PDP-11) byte order, interpreted as a
                UNIX date.

meldate   A four-byte value in middle-endian (PDP-11) byte order, interpreted as a
                UNIX-style date, but interpreted as local time rather than UTC.

regex   A regular expression match in extended POSIX regular expression syntax
                (like egrep). Regular expressions can take exponential time to process, and
                their performance is hard to predict, so their use is discouraged. When
                used in production environments, their performance should be carefully
                checked. The type specification can be optionally followed by /[c][s]. The
                "c" flag makes the match case insensitive, while the "s" flag update the
                offset to the start offset of the match, rather than the end. The regular
                expression is tested against line $N + 1$ onwards, where N is the given
                offset. Line endings are assumed to be in the machine's native format. ^
                and $ match the beginning and end of individual lines, respectively, not
                beginning and end of file.

search   A literal string search starting at the given offset. The same modifier flags
                can be used as for string patterns. The modifier flags (if any) must be
                followed by /number the range, that is, the number of positions at which
                the match will be attempted, starting from the start offset. This is suitable
                for searching larger binary expressions with variable offsets, using \
                escapes for special characters. The offset works as for regex.

default   This is intended to be used with the test *x* (which is always true) and a
                message that is to be used if there are no other matches.

clear   This test is always true and clears the match flag for that level. It is
                intended to be used with the default test.

name   Define a named magic instance that can be called from another use magic
                entry, like a subroutine call. Named instance direct magic offsets are
                relative to the offset of the previous matched entry, but indirect offsets are
                relative to the beginning of the file as usual. Named magic entries always
                match.

use   Recursively call the named magic starting from the current offset. If the
                name of the referenced instance begins with a ^ then the endianness of the
                magic is switched; if the magic mentioned leshort for example, it is treated
                as beshort and vice versa. This is useful to avoid duplicating the rules for
                different endianness.

Each top-level magic pattern (see below for an explanation of levels) is classified as text
or binary according to the types used. Types "regex" and "search" are classified as text
tests, unless non-printable characters are used in the pattern. All other tests are classified
as binary. A top-level pattern is considered to be a test text when all its patterns are text
patterns; otherwise, it is considered to be a binary pattern. When matching a file, binary
patterns are tried first; if no match is found, and the file looks like text, then its encoding is
determined and the text patterns are tried.

The numeric types may optionally be followed by & and a numeric value, to specify that
the value is to be AND'ed with the numeric value before any comparisons are done.
Prepending a u to the type indicates that ordered comparisons should be unsigned.

test   The value to be compared with the value from the file. If the type is numeric, this value is
                specified in C form; if it is a string, it is specified as a C string with the usual escapes
                permitted (e.g. \n for new-line).

Numeric values may be preceded by a character indicating the operation to be performed. It may be =, to specify that the value from the file must equal the specified value, <, to specify that the value from the file must be less than the specified value, >, to specify that the value from the file must be greater than the specified value, &, to specify that the value from the file must have set all of the bits that are set in the specified value, ^, to specify that the value from the file must have clear any of the bits that are set in the specified value, or ~, the value specified after is negated before tested. x, to specify that any value will match. If the character is omitted, it is assumed to be =. Operators &, ^, and ~ don't work with floats and doubles. The operator ! specifies that the line matches if the test does *not* succeed.

Numeric values are specified in C form; e.g. 13 is decimal, 013 is octal, and 0x13 is hexadecimal.

For string values, the string from the file must match the specified string. The operators =, < and > (but not &) can be applied to strings. The length used for matching is that of the string argument in the magic file. This means that a line can match any non-empty string (usually used to then print the string), with >\0 (because all non-empty strings are greater than the empty string).

The special test *x* always evaluates to true.

message   The message to be printed if the comparison succeeds. If the string contains a printf(3) format specification, the value from the file (with any specified masking performed) is printed using the message as the format string. If the string begins with "\b", the message printed is the remainder of the string with no whitespace added before it: multiple matches are normally separated by a single space.

A MIME type is given on a separate line, which must be the next non-blank or comment line after the magic line that identifies the file type, and has the following format:

   !:mime   MIMETYPE

i.e. the literal string "!:mime" followed by the MIME type.

Some file formats contain additional information which is to be printed along with the file type or need additional tests to determine the true file type. These additional tests are introduced by one or more > characters preceding the offset. The number of > on the line indicates the level of the test; a line with no > at the beginning is considered to be at level 0. Tests are arranged in a tree-like hierarchy: If a test on a line at level *n* succeeds, all following tests at level *n+1* are performed, and the messages printed if the

tests succeed, until a line with level *n* (or less) appears. For more complex files, one can use empty messages to get just the "if/then" effect, in the following way:

    0   string   MZ
    >0x18 leshort <0x40   MS-DOS executable
    >0x18 leshort >0x3f   extended PC executable (e.g., MS Windows)

Offsets do not need to be constant, but can also be read from the file being examined. If the first character following the last > is a *(* then the string after the parenthesis is interpreted as an indirect offset. That means that the number after the parenthesis is used as an offset in the file. The value at that offset is read, and is used again as an offset in the file. Indirect offsets are of the form: *((x [.[bslBSL]][+-][y]).* The value of *x* is used as an offset in the file. A byte, short or long is read at that offset depending on the [bslBSLm] type specifier. The capitalized types interpret the number as a big endian value, whereas the small letter versions interpret the number as a little endian value; the *m* type interprets the number as a middle endian (PDP-11) value. To that number the value of *y* is added and the result is used as an offset in the file. The default type if one is not specified is long.

That way variable length structures can be examined:

    # MS Windows executables are also valid MS-DOS executables
    0   string MZ
    >0x18   leshort <0x40   MZ executable (MS-DOS)
    # skip the whole block below if it is not an extended executable
    >0x18   leshort >0x3f
    >>(0x3c.l) string PE\0\0 PE executable (MS-Windows)
    >>(0x3c.l) string LX\0\0 LX executable (OS/2)

This strategy of examining has a drawback: You must make sure that you eventually print something, or users may get empty output (like, when there is neither PE\0\0 nor LE\0\0 in the above example)

If this indirect offset cannot be used directly, simple calculations are possible: appending *[+-*/%&/^]number* inside parentheses allows one to modify the value read from the file before it is used as an offset:

    # MS Windows executables are also valid MS-DOS executables
    0   string MZ
    # sometimes, the value at 0x18 is less that 0x40 but there's still an
    # extended executable, simply appended to the file
    >0x18   leshort <0x40
    >>(4.s*512) leshort 0x014c COFF executable (MS-DOS, DJGPP)

>>(4.s*512) leshort !0x014c MZ executable (MS-DOS)

Sometimes you do not know the exact offset as this depends on the length or position (when indirection was used before) of preceding fields. You can specify an offset relative to the end of the last up-level field using '&' as a prefix to the offset:

```
0        string  MZ
>0x18      leshort >0x3f
>>(0x3c.l)  string  PE\0\0    PE executable (MS-Windows)
# immediately following the PE signature is the CPU type
>>>&0       leshort 0x14c     for Intel 80386
>>>&0       leshort 0x184     for DEC Alpha
```

Indirect and relative offsets can be combined:

```
0          string  MZ
>0x18        leshort <0x40
>>(4.s*512)   leshort !0x014c MZ executable (MS-DOS)
# if it's not COFF, go back 512 bytes and add the offset taken
# from byte 2/3, which is yet another way of finding the start
# of the extended executable
>>>&(2.s-514) string  LE     LE executable (MS Windows VxD driver)
```

Or the other way around:

```
0            string  MZ
>0x18          leshort >0x3f
>>(0x3c.l)       string  LE\0\0  LE executable (MS-Windows)
# at offset 0x80 (-4, since relative offsets start at the end
# of the up-level match) inside the LE header, we find the absolute
# offset to the code area, where we look for a specific signature
>>>(&0x7c.l+0x26) string  UPX     \b, UPX compressed
```

Or even both!

```
0            string  MZ
>0x18          leshort >0x3f
>>(0x3c.l)      string  LE\0\0 LE executable (MS-Windows)
# at offset 0x58 inside the LE header, we find the relative offset
# to a data area where we look for a specific signature
```

>>>&(&0x54.l-3) string UNACE \b, ACE self-extracting archive

Finally, if you have to deal with offset/length pairs in your file, even the second value in a parenthesized expression can be taken from the file itself, using another set of parentheses. Note that this additional indirect offset is always relative to the start of the main indirect offset.

```
0              string     MZ
>0x18          leshort    >0x3f
>>(0x3c.l)      string      PE\0\0 PE executable (MS-Windows)
# search for the PE section called ".idata"...
>>>&0xf4          search/0x140 .idata
# ...and go to the end of it, calculated from start+length;
# these are located 14 and 10 bytes after the section name
>>>>(&0xe.l+(-4)) string      PK\3\4 \b, ZIP self-extracting archive
```

SEE ALSO
file(1) - the command that reads this file.

BUGS
The formats long, belong, lelong, melong, short, beshort, leshort, date, bedate, medate, ledate, beldate, leldate, and meldate are system-dependent; perhaps they should be specified as a number of bytes (2B, 4B, etc), since the files being recognized typically come from a system on which the lengths are invariant.